

## Calculations on tables.

### Goods Table

We have a goods table with 24,00,000 records

size of 1 goods entry=54 bytes.

Blocking factor=floor(512/54)=9. => 9 records per block.

So no. of blocks required=ceiling(24,00,000/9)=266667 blocks.

#### Case 1:

No ordering based on any field.

Linear search will take  $O(n/2)$  block accesses,

so approximately  $266667/2 = 133334$  block accesses.

#### Case 2:

Ordering based on goodsid,

so binary search would require ceiling(log(266667) to the base 2)=19 block accesses.

#### Case 3:

Index based on goodsid,

Size of goodsid=4 bytes, block ptr size = 6 bytes.

So size of index entry=4+6=10 bytes.

So no of index entries per block=floor(512/10)=51 entries.

Total no. of index entries = total no. of blocks in data file = 266667

so index blocking factor(no. of blocks required to store index entries,bfri)=  
ceiling( 266667/51)=5229 blocks

so now search will take ceiling(log(5229) to base 2)+1=13+1=14 block accesses

#### Case 4:

Multilevel index on goodsid,

1<sup>st</sup> level index(bfri) = 5229 entries

2<sup>nd</sup> level index = ceiling(5229/51) = 103 entries

3<sup>rd</sup> level index = ceiling(103/51)=3 entries

4<sup>th</sup> level index = ceiling(3/51)=1 entry.

So no. of block accesses = 4(level of indexing) + 1=5 block accesses.

#### Case 5: (for search based on brand)

Clustered index on brand, ordered by brand

We have 24,000 brands each having 100 goods.

Size of brand=10 bytes

Size of block ptr = 6 bytes.

Size of index entry = 10+6=16 bytes.

So no. of index entries per block =  $\text{floor}(512/16) = 32$  entries.

no. of index entries = no. of brands = 24,000

so no. of blocks =  $\text{ceiling}(24,000/32) = 750$  blocks.

So no. of block accesses by binary search =  $\text{ceiling}(\log(750) \text{ to base } 2) + 1 = 11$  block accesses

#### **Case 6:**

Multilevel clustered index on brand, assuming ordered by brand

Size of brand = 10 bytes

Size of block ptr = 6 bytes.

Size of index entry =  $10 + 6 = 16$  bytes.

So no. of index entries per block =  $\text{floor}(512/16) = 32$  entries.

no. of index entries = no. of brands = 24,000

so no. of blocks =  $\text{ceiling}(24,000/32) = 750$  blocks.

=>

1<sup>st</sup> level index = 750

2<sup>nd</sup> level index =  $\text{ceiling}(750/32) = 24$

3<sup>rd</sup> level index =  $\text{ceiling}(24/32) = 1$

so no. of block accesses =  $3(\text{level of indexing}) + 1 = 4$  block accesses.

#### **Case 7:**

B-Tree Index on goodsid

Block pointer size P = 6 bytes

Record pointer size R = 7 bytes

Size of search key field = 4 bytes

Block size B = 512 bytes

Number of goods = 24 lakh

A node in a B-Tree can have at most p pointers, p-1 entries and p-1 data pointers. Since all these need to fit within a block, we have :

$$6p + 4(p-1) + 7(p-1) \leq 512$$

p approximately equal to 29

Assuming that 69 percent of the nodes are full when the number of values in the tree stabilizes, we have  $p = 0.69p = 20$

No. of levels =  $\log(2400000)/\log(20) = 5$  levels

If we don't consider the root we have levels = 4

Level	Nodes	Entries	Pointers
Root	1	19	20
Level 1	20	380	400
Level 2	400	7600	8000
Level 3	8000	1,52,000	1,60,000
Level 4	1,60,000	30,40,000	

No. of entries and pointers barring the last level = 328419

Remaining entries = 2400000-159999 = 2071581

No. of nodes in the last level = Remaining entries+pointers/no. of entries+pointers in one node  
=53118

Since size of one node <= size of a block, we will have to fit one node per block

Therefore number of blocks = 53118

Number of block accesses = level +1= 5 (not including root)

Summarizing.

Property	Block Accesses	No. of extra blocks required
No ordering	133334	0
Ordering on goods id	20	0
Index on goods name	14	5229
Multilevel index on goods id	5	5229+103+3+1=5363
Clustered Index on brand	11	750
Multilevel Clustered Index on brand	4	750+24+1=775
B-tree index on goods id	5	53118+21+400+8000 = 61538

Looking at these figures, we conclude that using a Multilevel Index on goods id is better than B-Tree Index on goods id.

We will also include Multilevel clustered index on brand to optimize our search queries.

## Seller Table

We have a seller table with 10,000 records.

Seller record size=48 bytes.

Blocking factor=floor(512/48)=10 entries per block

No. of blocks required = ceiling(10,000/10)=1000

### Case 1:

No ordering based on any field.

Linear search will take  $O(n/2)$  block accesses,  
so approximately  $1000/2 = 500$  block accesses.

### Case 2:

Ordering based on sid,

so binary search would require ceiling(log(1000) to the base 2)=10 block accesses.

### Case 3:

Index based on sid,

Size of sid=20 bytes, block ptr size = 6 bytes.

So size of index entry=20+6=26 bytes.

So no. of index entries per block=floor(512/26)=19 entries.

Total no. of index entries = total no. of blocks in data file = 1000

so index blocking factor(no. of blocks required to store index entries,bfri)=  
ceiling( 1000/19)=53 entries

so now search will take ceiling(log(53) to base 2)=6 block accesses

### Case 4:

Multilevel index on sid,

1<sup>st</sup> level index(bfri) = 53 entries

2<sup>nd</sup> level index = ceiling(53/19) = 3 entries

3<sup>rd</sup> level index = ceiling(3/19)=1 entry

So no. of block accesses = 3(level of indexing) + 1=4 block accesses.

### Case 5:

Btree : Index on Seller id

Block pointer size P = 6 bytes

Record pointer size R = 7 bytes

Size of search key field = 20 bytes

Block size B = 512 bytes

Number of goods = 10000

A node in a B-Tree can have at most  $p$  pointers,  $p-1$  entries and  $p-1$  data pointers. Since all these need to fit within a block, we have :

$$6p + 20(p-1) + 7(p-1) \leq 512$$

$p$  approximately equal to 15

Assuming that 69 percent of the nodes are full when the number of values in the tree stabilizes, we have  $p = 0.69p = 11$

No. of levels =  $\log(10000)/\log(11) = 4$  levels  
 If we don't consider root levels = 3

Level	Nodes	Entries	Pointers
Root	1	10	11
Level 1	11	110	121
Level 2	121	1210	1331
Level 3	1331	13310	14641

No. of entries and pointers barring the last level = 2793  
 Remaining entries =  $10000 - 2793 = 7207$

No. of nodes in the last level = Remaining entries+pointers/no. of entries+pointers in one node =  $7207/21 = 343$   
 Since size of one node  $\leq$  size of a block, we will have to fit one node per block

Therefore number of blocks = 343  
 Number of block accesses = level +1 = 4

Summarizing.

Property	Block accesses	No. of extra blocks required
No ordering	500	0
Ordering on seller id	11	0
Index on seller id	7	53
Multilevel index on seller id	4	$53+3+1=57$
B-tree index on seller id	4	$343+121+11+1=476$

We have chose Multilevel index on Seller Id

## Customer Table

We have a customer table with 24,000 records.

Customer record size=154 bytes.

Blocking factor= $\text{floor}(512/154) = 3$  entries per block.

No. of blocks required =  $\text{ceiling}(24,000/3) = 8000$ .

### Case 1:

No ordering based on any field.

Linear search will take  $O(n/2)$  block accesses,  
so approximately  $8000/2 = 4000$  block accesses.

### Case 2:

Ordering based on custid,

so binary search would require  $\text{ceiling}(\log_2(8000)) = 13$  block accesses.

### Case 3:

Index based on custid,

Size of custid=20 bytes, block ptr size = 6 bytes.

So size of index entry=20+6=26 bytes.

So no. of index entries per block= $\text{floor}(512/26)=19$  entries.

Total no. of index entries = total no. of blocks in data file = 8000

so index blocking factor(no. of blocks required to store index entries,bfri)=  
 $\text{ceiling}(8000/19)=422$  entries

so now search will take  $\text{ceiling}(\log_2(422)) + 1 = 9 + 1 = 10$  block accesses.

### Case 4:

Multilevel index on goodsid,

1<sup>st</sup> level index(bfri) = 422 entries

2<sup>nd</sup> level index =  $\text{ceiling}(422/19) = 23$  entries

3<sup>rd</sup> level index =  $\text{ceiling}(23/19) = 2$  entries

4<sup>th</sup> level index =  $\text{ceiling}(2/19) = 1$  entry.

So no. of block accesses = 4(level of indexing) + 1 = 5 block accesses.

### Case 5:

Btree : Index on Customer id

Block pointer size P = 6 bytes

Record pointer size R = 7 bytes

Size of search key field = 20 bytes

Block size B = 512 bytes

Number of goods = 24000

A node in a B-Tree can have at most  $p$  pointers,  $p-1$  entries and  $p-1$  data pointers. Since all these need to fit within a block, we have :

$$6p + 20(p-1) + 7(p-1) \leq 512$$

$p$  approximately equal to 15

Assuming that 69 percent of the nodes are full when the number of values in the tree stabilizes, we have  $p = 0.69p = 11$

No. of levels =  $\log(24000)/\log(11) = 5$  levels  
 If we don't consider the root level = 4

Level	Nodes	Entries	Pointers
Root	1	10	11
Level 1	11	110	121
Level 2	121	1210	1331
Level 3	1331	13310	14641
Level 4	14641	146410	161051

No. of entries and pointers barring the last level = 2793  
 Remaining entries =  $24000 - 2793 = 21207$

No. of nodes in the last level = Remaining entries and pointers/no. of entries and pointers in one node = 1009  
 Since size of one node  $\leq$  size of a block, we will have to fit one node per block

Therefore number of blocks = 1009  
 Number of block accesses = level + 1 = 5  
 Summarizing.

Property	Block accesses	No. of extra blocks required
No ordering	4000	0
Ordering on customer id	14	0
Index on customer id	10	422
Multilevel index on customer id	5	$422 + 23 + 2 + 1 = 448$
B-Tree Index on customer id	5	$1009 + 1 + 11 + 121 = 1142$

We have decided to use Multilevel on Customer Id

## Products Table

We have a products table with 36,00,000 records  
Customer record size=154 bytes.116

Blocking factor=floor(512/116)= 4 entries per block.

No. of blocks required = ceiling(3600000/4)= 900000.

### Case 1:

No ordering based on any field.

Linear search will take  $O(n/2)$  block accesses,  
so approximately  $900000/2 = 450000$  block accesses.

### Case 2:

Ordering based on pid,

so binary search would require ceiling(log(450000) to the base 2) =23 block accesses.

### Case 3:

Index based on pid,

Size of pid=20 bytes, block ptr size = 6 bytes.

So size of index entry=20+6=26 bytes.

So no. of index entries per block=floor(512/26)=19 entries.

Total no. of index entries = total no. of blocks in data file = 3600000

so index blocking factor(no. of blocks required to store index entries,bfri)=  
ceiling( 3600000/19)=189474 entries

so now search will take ceiling(log(189474) to base 2)+1=18+1=19 block accesses.

### Case 4:

Multilevel index on pid,

1<sup>st</sup> level index(bfri) = 189474 entries

2<sup>nd</sup> level index = ceiling(189474/19) = 9973entries

3<sup>rd</sup> level index = ceiling(9973/19)=525 entries

4<sup>th</sup> level index = ceiling(525/19)=28 entries.

5<sup>th</sup> level index = ceiling(28/19)=2 entries

6<sup>th</sup> level index = ceiling(2/19)=1 entry

So no. of block accesses = 6(level of indexing) + 1=7 block accesses.

### Case 5:

Btree : Index on Product id

Block pointer size P = 6 bytes

Record pointer size R = 7 bytes

Size of search key field = 20 bytes

Block size B = 512 bytes



Number of goods = 3600000

A node in a B-Tree can have at most p pointers, p-1 entries and p-1 data pointers. Since all these need to fit within a block, we have :

$$6p + 20(p-1) + 7(p-1) \leq 512$$

p approximately equal to 15

Assuming that 69 percent of the nodes are full when the number of values in the tree stabilizes, we have  $p = 0.69p = 11$

$$\text{No. of levels} = \log(3600000)/\log(11) = 7 \text{ levels}$$

If we don't consider the root then levels = 6

Level	Nodes	Entries	Pointers
Root	1	10	11
Level 1	11	110	121
Level 2	121	1210	1331
Level 3	1331	13310	14641
Level 4	14641	146410	161051
Level 5	161051	1610510	1771561
Level 6	1771561	17715610	

No. of entries and pointers barring the last level = 338205

Remaining entries = 3600000-338205 = 3261795

No. of nodes in the last level = Remaining entries+pointers/no. of entries+pointers in one node = 155324

Since size of one node  $\leq$  size of a block, we will have to fit one node per block

Therefore number of blocks = 155324

Number of block accesses = level + 1 = 7

Summarizing.

Property	Block Accesses	No. of extra blocks required
No ordering	450000	0
Ordering on pid	24	0
Index on pid	19	189474
Multilevel index on pid	7	$189474+525+28+2+1=190030$
Btree index on pid	7	$155324+1+11+121+1331+14641=171429$

Since the number of block accesses are lesser in the case of BTree index on pid as opposed to Multilevel index on pid, we have chosen this.

## Books Table

We have 600000 records

Index entry = 4bytes + 6 bytes [from goodsid + blockptr]

$$Bfri = 512/10 = 51$$

So we need  $\lceil 600000/51 \rceil = 11765$  blocks

$$\text{So block access} = \log(11765)/\log(2) + 1 = 15$$

## Media Table

We have 600000 records

Index entry = 4bytes + 6 bytes [from goodsid + blockptr]

$$Bfri = 512/10 = 51$$

So we need  $\lceil 600000/51 \rceil = 11765$  blocks

$$\text{So block access} = \log(11765)/\log(2) + 1 = 15$$

## Fashion Table

We have 600000 records

Index entry = 4bytes + 6 bytes [from goodsid + blockptr]

$$Bfri = 512/10 = 51$$

So we need  $\lceil 600000/51 \rceil = 11765$  blocks

$$\text{So block access} = \log(11765)/\log(2) + 1 = 15$$

## TV Table

We have 600000 records

Index entry = 4bytes + 6 bytes [from goodsid + blockptr]

$$Bfri = 512/10 = 51$$

So we need  $\lceil 600000/51 \rceil = 11765$  blocks

$$\text{So block access} = \log(11765)/\log(2) + 1 = 15$$

## Laptop Table

We have 600000 records

Index entry = 4bytes + 6 bytes [from goodsid + blockptr]

$$Bfri = 512/10 = 51$$

So we need  $\lceil 600000/51 \rceil = 11765$  blocks

$$\text{So block access} = \log(11765)/\log(2) + 1 = 15$$

## Mobiles Table

We have 600000 records

Index entry = 4bytes + 6 bytes [from goodsid + blockptr]

$B_{fri} = 512/10 = 51$

So we need  $\lceil 600000/51 \rceil = 11765$  blocks

So block access =  $\log(11765)/\log(2) + 1 = 15$

## QUERY EXECUTION

Assuming 1 Block access time=24.11ms

1)select \* from goods where goodsid="x" or brand="x"

This is equivalent to saying the no. of block accesses for

(select \* from goods where goodsid="x")

+

(select \* from goods where brand="x")

- since we have a multilevel index on goodsid on goodsid in goods, first query will take 5 block accesses
- since we have a multilevel clustered index on brand in goodsid, second query will take 4 block accesses

so the whole query will take 9 block accesses.

So avg. query execution time =  $9 * 24.11 = 217\text{ms}$

2)select \* from products join goods on products.goodsid=goods.goodsid where goodsid="x";

- since we have a multilevel index on goodsid in goods table(5),
- since we have a btree index on pid in products table(7),

the query will take  $5 * 7$  block accesses = 35 block accesses

So avg. query execution time =  $35 * 24.11 = 843.85\text{ms}$

3)select \* from seller join products on seller.sid=products.sid where sid=x;

- since we have a multilevel index on sid in seller table(4 block accesses),
- since we have no index on sid in products table( $\log(3600000)$  to base 2 block accesses, as we have 3600000 records in products table),

the query will take  $22 \times 4$  block accesses = 88 block accesses

So avg. query execution time =  $88 \times 24.11 = 2121\text{ms}$

4)select \* from products where pid="x";

- since we have a btree index on pid in products table we'll need 7 block accesses.

So avg. query execution time =  $7 \times 24.11 = 169\text{ms}$

5)select \* from customers where customer="x"

- since we have a multilevel index on custid in customers table we'll need 5 block accesses.

So avg. query execution time =  $5 \times 24.11 = 121\text{ms}$

